

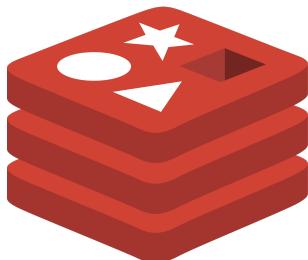
# Redis

CST 363

# What is Redis?

## REmote DIctionary Server

- An open-source, in-memory, single-threaded "data structure store"
  - allows for the storage and retrieval of data structures, such as strings, hashes, lists, sets, and more
  - speedy – micro-seconds vs. the milli-seconds we're used to with disk-based stores.
- Known for its high performance, scalability, and versatility
- Virtually every production system uses Redis (or an equivalent) for caching, sessions, rate-limiting, real-time messaging, and lightweight queues.
  - A staple of "database adjacencies."



# History

- 2009 - Salvatore Sanfilippo, an Italian developer, found that MySQL could not provide the necessary performance for real-time web analytics
- Went from a small personal project from Sanfilippo, to being an industry standard for in-memory data storage.



# Comparison to SQL

Aspect	SQL Databases	Redis
Data Model	Relational model with structured tables	NoSQL key-value store with support for multiple data structures (e.g., lists, sets, hashes, streams)
Storage	Disk-based storage with caching layers for performance	In-memory storage for speed, with optional disk persistence (RDB, AOF)
Query Language	Uses SQL (Structured Query Language)	Uses its own command set (e.g., GET, SET, HGETALL); no SQL
Use Cases	Structured data, complex queries, transactional integrity, relational models	Caching, real-time analytics, messaging, fast-access scenarios

# Where Traditional RDBMSs Struggle (1 of 5)

## Hot-Read Caching

- A small subset of data ("hot keys" or documents) is read far more often (10x - 100x) than the rest.
- Every cache-miss triggers planner, optimizer, disk I/O (~1-3 ms each)

## With Redis:

- `GET user:1234` always sub-millisecond
- Built-in "Least Recently Used" (LRU) eviction, no extra query planning

# Where Traditional RDBMSs Struggle (2 of 5)

## High-Frequency Counters & Concurrency

- Atomic `UPDATE ... SET counter = counter + 1` at very high queries per second (QPS)
- Row-level locks slow things down under heavy load.
- WAL/journaling thrash + autovacuum bloat slow sustained writes

## With Redis:

- `INCR page:view:5678` is lock-free &  $O(1)$

# Where Traditional RDBMSs Struggle (3 of 5)

## Ephemeral State & Session Management

- Short-lived data (sessions, feature-flags, temp locks) that expires on its own.
- Frequent `INSERT / DELETE` churn, index maintenance
- Table bloat → expensive `VACUUM`, fragmentation

## With Redis:

- Keys with `TTL` auto-expire; no background `GC`
- In-memory list, set, hash structures for fast lookups

# Where Traditional RDBMSs Struggle (4 of 5)

## Event Streaming & Pub/Sub

- On-disk queues or tailable cursors lack true ordering & replay
- Notifications don't scale to many consumers
- No built-in consumer groups or at-least-once delivery guarantees

**With Redis:** Redis Streams + consumer-groups for ordered, replayable log

# Where Traditional RDBMSs Struggle (5 of 5)

## **Heavy Analytical & Batch Workloads**

- Large scans, complex joins or aggregations, real-time analytics.
- Long-running queries lock resources, contend with OLTP traffic, and can overwhelm primary nodes.

## **With Redis:**

- Can use RedisTimeSeries or Lua scripts for near-real-time stats

# Core Data Structures

Redis supports five different data structures: strings, hashes, lists, sets and ordered sets – regardless of the type, a value is accessed by a key.

- **String**: arbitrary byte values
- **Hash**: maps of fields → values
- **List**: ordered collections
- **Set**: unique, unordered collections
- **Sorted Set**: scored, ordered collections

Data type	Canonical commands	Typical use
<b>String</b>	SET/GET/INCRBY	Counters, flags
<b>Hash</b>	HSET/HGETALL	User profiles
<b>List</b>	RPUSH/LRANGE	Log buffers, task queues
<b>Set</b>	SADD/SMEMBERS	Unique tags
<b>Sorted Set</b>	ZADD/ZRANGE	Leaderboards, feeds
<b>Streams</b>	XADD/XREAD	Append-only event logs

# Memory-Oriented Caveats (1 / 2)

Redis is fast **because** everything lives in RAM—here's how to keep it that way.

Lever	Why it matters	Quick tips
<code>maxmemory</code>	Hard ceiling for RAM use	Size for worst-case + headroom • Monitor with <code>INFO MEMORY</code>
<b>Eviction policy</b>	Decides which keys disappear when full	<code>allkeys-lru</code> (smart cache) • <code>volatile-ttl</code> (expire-only) • Benchmark with <code>redis-benchmark --lru</code>
<b>Big-key anti-pattern</b>	5 MB key blocks event loop, slows replicas & AOF rewrite	Shard large hashes/lists • Prefer many small keys

# Memory-Oriented Caveats (2 / 2)

Lever	Why it matters	Quick tips
<b>Memory-efficient structures</b>	Same task, less RAM	Bitmaps for flags • HyperLogLog for cardinality • Bloom filters for “probable existence”
<b>Diagnostics</b>	Spot trouble early	<code>MEMORY USAGE &lt;key&gt;</code> • <code>MEMORY DOCTOR</code> • <code>MEMORY STATS</code>

**Rule of thumb:** If a key can't be read in <1 ms or fits on one screen, it's probably too big—break it up or pick a leaner structure.

# Hands-on: Basic Commands

Access the Redis command line interface (redis-cli) with:

```
docker exec -it my-redis redis-cli
```

## 1. Strings & counters

```
SET page_views 0
INCRBY page_views 42
GET page_views # 42
```

## 2. Lists

```
RPUSH recent_log "user123 signup" "user123 click"
LRANGE recent_log 0 -1
```

# Hands-on: Basic Commands (cont.)

## 3. Sets & membership

```
SADD online_users u1 u2 u3  
SCARD online_users  
SISMEMBER online_users u2
```

## 4. Pub/Sub teaser

```
SUBSCRIBE classroom  
# in another shell  
PUBLISH classroom "Hello CST 363"
```

# Redis Persistence & Durability

"Isn't everything lost on reboot?"

Mechanism	What it does	Write frequency	Disk footprint	Worst-case data loss*
RDB snapshot	Forks the process and saves a compressed dump file ( .rdb )	On a schedule (e.g., <code>save 900 1 , save 300 10</code> )	Small, binary & compressed	All writes since last snapshot
AOF (Append-Only File)	Appends every write to a log and replays it on restart	<code>appendfsync always   everysec   no</code>	Larger, plaintext stream	0 s with <code>always</code> ; $\leq 1$ s with <code>everysec</code>

\*"Data loss" here means writes not yet persisted if the server crashes.

# Hybrid best-practice

```
appendonly yes          # enable AOF
appendfsync everysec   # fsync once per second
save 3600 1            # RDB snapshot every hour
```

*Blend near-zero data loss (AOF) with compact hourly snapshots (RDB) for quick restarts & cheaper off-box backups.*

# Redis + PostgreSQL ❤️

- **Read-through cache**
- Flask route hits Redis → miss triggers Postgres query → result cached with TTL.
- **Write-behind / event sourcing**
  - App writes to Redis **Stream**; separate worker persists to Postgres asynchronously.
- **Real-time counters**
  - `INCR` in Redis, nightly ETL to Postgres for durable analytics.
- **Pub/Sub fan-out** for web-socket notifications while Postgres stays source of truth.
  - Publisher sends one message to a channel
  - Redis instantly "fans out" (copies) that message to every active subscriber on that channel.

# NBA Play-by-Play App

Real-time streaming + durable storage

- "Pretend live" multi-game replay of NBA play-by-play
- Accelerated or real-time playback via `SIM_SPEED`
- **Redis** for low-latency fan-out & caching
- **Postgres** for durable event logging
- Browser dashboard via Flask + Server-Sent Events



# Data Ingestion & Scheduling

1. **Read CSV** ( `pbp2023.csv` ) into Pandas
2. **Filter** to a few games on a certain date (here 2023-03-24)
3. **Compute** `elapsed_seconds` per event:

```
elapsed = minutes*60 + seconds + periods_before*period_length
```

4. **Seed a min-heap** with each game's first event

```
heap = [(elapsed0, game_id, idx0, events0), ... ]
```

# Real-time Replay Loop

- **Pop** next event from heap → `(sim_t, gid, idx, evts)`
- **Sleep** until `(sim_t / SIM_SPEED)` matches wall-clock
- **Publish** payload → Redis pub/sub + cache
- **Batch** payload → Postgres `INSERT ... executemany`
- **Push** next event of that game back onto heap
- Repeat until heap is empty

# Redis: Real-time Layer

- **Pub/Sub**

- `r.publish("pbp_live", payload_json)`
- Flask SSE subscribes → pushes to browsers

- **Cache**

- `r.set("game:{id}:latest", payload_json, ex=120)`
- Quick lookup of current score/state

# Postgres: Durable Log

- Table `pbp(ts timestampz, game_id text, ...)`
- Batched writes ( 200 events per COMMIT)
- Index on `(game_id, ts)` for fast history queries
- Enables audit, replay, analytics after the demo

# Flask + SSE Front-End

- `/` → renders Bootstrap dashboard
- `/events` → SSE endpoint using `redis.pubsub().listen()`
- Browser JS:

```
const es = new EventSource("/events");
es.onmessage = ev => { ... }
```

- Updates:
  - Scoreboard cards per game
  - Live feed list (last 100 plays)

# Why This Architecture?

Concern	Tool	Role
Real-time updates	Redis	Pub/Sub & in-memory cache
Durable persistence	Postgres	Batched event store
Simple streaming	SSE	Server-Sent Events → browser EventSource
Merge multiple feeds	Heap	Time-ordered “live” event scheduling

# Possible Enhancements

- **Scale fan-out:** Redis Streams, Kafka, or WebSockets
- **Async I/O:** FastAPI/Quart + async Redis client
- **Historical API:** add `/history?game_id=` querying Postgres
- **Schema enhancements:** numeric time fields, PKs, constraints
- **Simplify demo:** flatten+sort or `heapq.merge` for static CSV

# Redis Docs

Check out the [Redis website](#) for documentation and tutorials!